

## ▼ Load and Combine Datasets

```
import numpy as np
import pandas as pd
```

The Census Profile contains over 2600 "characteristics" for each region, stored in a one-row-per-characteristic format, with columns representing various metadata and breakdowns by gender. In this section, we will select only those characteristics that correspond to our desired feature domains (age, housing, income, family & citizenship, education, employment, and commuting patterns). These are then normalized over the population/dwelling count, and transposed into the expected columnar format for further analysis as features. Characteristics in the Census Profile often have a hierarchical structure, and the provided names may not make sense outside of this. We will rename these variables in order to properly reflect their meaning.

```

# Assigning meaningful variable names to Census Profile characteristics in place of numeric IDs
all_vars = ["total_pop", "total_dwellings", "num_male",
            "age_25", "age_30", "age_35", "age_40", "age_45", "age_50", "age_55", "age_60", "age_65_up", "age_mean", "age_med",
            "homes_detached_house", "homes_semidetached_house", "homes_rowhouse", "homes_duplex_apt", "homes_lowrise_apt",
            "homes_highrise_apt", "homes_other_stationary", "homes_mobile",
            "avg_ppl_household", "married_ppl", "single_ppl",
            "avg_total_income", "avg_aftertax_income", "med_fulltime_income", "avg_fulltime_income",
            "income_none", "income_under_10k", "income_10k", "income_20k", "income_30k", "income_40k", "income_50k", "income_60k",
            "income_70k", "income_80k", "income_90k", "income_100k_up",
            "med_total_household_income", "med_aftertax_household_income",
            "indigenous_ppl", "nonindigenous_ppl",
            "home_owner", "home_renter", "home_gov_or_ind_band",
            "condo", "non_condo", "avg_rooms_home",
            "home_cost_under_30pct", "home_cost_over_30pct",
            "can_citizen_ppl", "non_citizen_ppl",
            "non_immigrant_ppl", "immigrant_ppl", "non_perm_res_ppl",
            "no_move_last_yr", "moved_last_yr", "no_move_5yrs", "moved_last_5yrs",
            "school_no_hs", "school_hs", "school_college", "school_uni_degree",
            "edu_field_education", "edu_field_arts_comms", "edu_field_humanities", "edu_field_socsci_law", "edu_field_bus_admin",
            "edu_field_science", "edu_field_math_cs", "edu_field_engin_arch", "edu_field_agri_res_env", "edu_field_health",
            "edu_field_pers_protect_transp", "edu_field_other",
            "workforce_participation_rate", "workforce_employment_rate", "workforce_unemployment_rate",
            "work_lastyear_didnotwork", "work_lastyear_worked", "work_lastyear_fulltime", "work_lastyear_parttime", "work_lastyear_selfemp",
            "worktype_employee", "worktype_selfemp",
            "occup_cat_snr_mgmt", "occup_cat_busfin", "occup_cat_science", "occup_cat_health", "occup_cat_edu_law_socserv",
            "occup_cat_sales_serv", "occup_cat_trades_transp", "occup_cat_natres_agr", "occup_cat_manuf_util",
            "occup_ind_agr_forest", "occup_ind_mine_og", "occup_ind_util", "occup_ind_constr", "occup_ind_manuf", "occup_ind_retail_trd",
            "occup_ind_transp_warehs", "occup_ind_info_culture", "occup_ind_fin_insure", "occup_ind_prof_sci_tech_serv",
            "occup_ind_mgmt", "occup_ind_admsupport_wastemgmt", "occup_ind_edu", "occup_ind_arts_ent_rec",
            "occup_ind_accom_food_svc", "occup_ind_other", "occup_ind_pubadmin",
            "work_loc_home", "work_loc_foreign", "work_loc_notfixed", "work_loc_workplace",
            "commute_same_subdiv", "commute_same_div", "commute_same_prov", "commute_diff_prov",
            "commute_transp_cardriver", "commute_transp_carpass", "commute_transp_pubtrans", "commute_transp_walk",
            "commute_transp_bike", "commute_transp_other",
            "commute_time_under15", "commute_time_15", "commute_time_30", "commute_time_45", "commute_time_over60",
            "commute_start_5am", "commute_start_6am", "commute_start_7am", "commute_start_8am", "commute_start_9am", "commute_start_10am", "commute_start_11am", "commute_start_12pm", "commute_start_1pm", "commute_start_2pm", "commute_start_3pm", "commute_start_4pm", "commute_start_5pm", "commute_start_6pm", "commute_start_7pm", "commute_start_8pm", "commute_start_9pm", "commute_start_10pm", "commute_start_11pm", "commute_start_12am"]

var_names = dict(zip(
    [1, 4, 8, 16, 17, 18, 19, 20, 21, 22, 23, 24, 39, 40,
     42, 43, 44, 45, 46, 47, 48, 49, 57, 59, 66, 128, 130, 143, 144,
     156, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168,
     243, 244, 1403, 1410, 1415, 1416, 1417, 1419, 1420, 1433,
     1466, 1467, 1523, 1526, 1528, 1529, 1537, 1975, 1976, 1984, 1985,
     2015, 2016, 2018, 2024, 2095, 2097, 2100, 2109, 2117, 2121, 2127, 2132, 2140, 2143,
     2149, 2155, 2228, 2229, 2230, 2232, 2233, 2234, 2235, 2236, 2240, 2245,
     2249, 2250, 2251, 2252, 2253, 2254, 2255, 2256, 2257, 2258,
     2262, 2263, 2264, 2265, 2266, 2267, 2268, 2269, 2270,
     2271, 2272, 2273, 2274, 2275, 2276, 2277, 2278, 2279, 2280, 2281,
     2594, 2595, 2596, 2597,
     2599, 2600, 2601, 2602,
     2605, 2606, 2607, 2608, 2609, 2610,
     2612, 2613, 2614, 2615, 2616,
     2618, 2619, 2620, 2621, 2622, 2623],
    all_vars
))

# Divide variables into those that represent counts of people and dwellings, in order to divide by total population
ppl_vars = ["num_male", "age_25", "age_30", "age_35", "age_40", "age_45", "age_50", "age_55", "age_60", "age_65_up",
            "married_ppl", "single_ppl", "income_none", "income_under_10k", "income_10k", "income_20k", "income_30k",
            "income_40k", "income_50k", "income_60k", "income_70k", "income_80k", "income_90k", "income_100k_up",
            "indigenous_ppl", "nonindigenous_ppl", "home_owner", "home_renter", "home_gov_or_ind_band",
            "home_cost_under_30pct", "home_cost_over_30pct", "can_citizen_ppl", "non_citizen_ppl", "non_immigrant_ppl",
            "immigrant_ppl", "non_perm_res_ppl", "no_move_last_yr", "moved_last_yr", "no_move_5yrs", "moved_last_5yrs",
            "school_no_hs", "school_hs", "school_college", "school_uni_degree", "edu_field_education", "edu_field_arts_comms",
            "edu_field_humanities", "edu_field_socsci_law", "edu_field_bus_admin", "edu_field_science", "edu_field_math_cs",
            "edu_field_engin_arch", "edu_field_agri_res_env", "edu_field_health", "edu_field_pers_protect_transp", "edu_field_other",
            "workforce_participation_rate", "workforce_employment_rate", "workforce_unemployment_rate",
            "work_lastyear_didnotwork", "work_lastyear_worked", "work_lastyear_fulltime", "work_lastyear_parttime", "work_lastyear_selfemp",
            "worktype_employee", "worktype_selfemp",
            "occup_cat_snr_mgmt", "occup_cat_busfin", "occup_cat_science", "occup_cat_health", "occup_cat_edu_law_socserv",
            "occup_cat_sales_serv", "occup_cat_trades_transp", "occup_cat_natres_agr", "occup_cat_manuf_util",
            "occup_ind_agr_forest", "occup_ind_mine_og", "occup_ind_util", "occup_ind_constr", "occup_ind_manuf", "occup_ind_retail_trd",
            "occup_ind_transp_warehs", "occup_ind_info_culture", "occup_ind_fin_insure", "occup_ind_prof_sci_tech_serv",
            "occup_ind_mgmt", "occup_ind_admsupport_wastemgmt", "occup_ind_edu", "occup_ind_arts_ent_rec",
            "occup_ind_accom_food_svc", "occup_ind_other", "occup_ind_pubadmin",
            "work_loc_home", "work_loc_foreign", "work_loc_notfixed", "work_loc_workplace",
            "commute_same_subdiv", "commute_same_div", "commute_same_prov", "commute_diff_prov",
            "commute_transp_cardriver", "commute_transp_carpass", "commute_transp_pubtrans", "commute_transp_walk",
            "commute_transp_bike", "commute_transp_other",
            "commute_time_under15", "commute_time_15", "commute_time_30", "commute_time_45", "commute_time_over60",
            "commute_start_5am", "commute_start_6am", "commute_start_7am", "commute_start_8am", "commute_start_9am", "commute_start_10am", "commute_start_11am", "commute_start_12pm", "commute_start_1pm", "commute_start_2pm", "commute_start_3pm", "commute_start_4pm", "commute_start_5pm", "commute_start_6pm", "commute_start_7pm", "commute_start_8pm", "commute_start_9pm", "commute_start_10pm", "commute_start_11pm", "commute_start_12am"]

```

```

"edu_field_engin_arch", "edu_field_agri_res_env", "edu_field_health", "edu_field_pers_protect_tr:
"work_lastyear_didnotwork", "work_lastyear_worked", "work_lastyear_fulltime", "work_lastyear_par:
"worktype_employee", "worktype_selfemp", "occup_cat_snr_mgmt", "occup_cat_busfin", "occup_cat_sci:
"occup_cat_edu_law_socserv", "occup_cat_arts_rec", "occup_cat_sales_serv", "occup_cat_trades_tra:
"occup_cat_manuf_util", "occup_ind_agr_forest", "occup_ind_mine_og", "occup_ind_util", "occup_ind:
"occup_ind_wholesale_trd", "occup_ind_retail_trd", "occup_ind_transp_warehs", "occup_ind_info_cu:
"occup_ind_realestate", "occup_ind_prof_sci_tech_serv", "occup_ind_mgmt", "occup_ind_admsupport_v:
"occup_ind_health_socasst", "occup_ind_arts_ent_rec", "occup_ind_accom_food_svc", "occup_ind_othe:
"work_loc_home", "work_loc_foreign", "work_loc_notfixed", "work_loc_workplace", "commute_same_sub:
"commute_same_prov", "commute_diff_prov", "commute_transp_cardriver", "commute_transp_carpass:
"commute_transp_walk", "commute_transp_bike", "commute_transp_other", "commute_time_under15", "co:
"commute_time_30", "commute_time_45", "commute_time_over60", "commute_start_5am", "commute_st:
"commute_start_8am", "commute_start_9am", "commute_start_noon"]

dwelling_vars = ["homes_detached_house", "homes_semidetached_house", "homes_rowhouse", "homes_duplex_apt", "homes_lowri:
                "homes_highrise_apt", "homes_other_stationary", "homes_mobile", "condo", "non_condo"]

```

```
# Load Census Profile data
```

```
census_data = pd.read_csv('98-401-X2021013_English_CSV_data.csv', encoding="iso-8859-1")
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
```

```
<ipython-input-4-364ce2b10ad5> in <cell line: 2>()
```

```
1 # Load Census Profile data
```

```
----> 2 census_data = pd.read_csv('98-401-X2021013_English_CSV_data.csv', encoding="iso-8859-1")
```

⏏ 6 frames

```
/usr/local/lib/python3.10/dist-packages/pandas/io/common.py in get_handle(path_or_buf, mode,
encoding, compression, memory_map, is_text, errors, storage_options)
```

```
854         if ioargs.encoding and "b" not in ioargs.mode:
```

```
855             # Encoding
```

```
--> 856         handle = open(
```

```
857             handle,
```

```
858             ioargs.mode,
```

```
FileNotFoundError: [Errno 2] No such file or directory: '98-401-
```

```
# Filter to Ontario FSAs/postal codes
```

```
census_data = census_data[census_data['GEO_NAME'].str.match(r'K|L|M|N|P')]
```

```
# Drop extra columns
```

```
census_data = census_data[["GEO_NAME", "CHARACTERISTIC_ID", "C1_COUNT_TOTAL", "C2_COUNT_MEN+"]]
```

```
# Use the number of men for the "num_male" value, otherwise just take the total
```

```
census_data['value'] = census_data.apply(lambda row: row['C2_COUNT_MEN+'] if row['CHARACTERISTIC_ID']==8 else
```

```
# Select only those rows corresponding to the characteristics we want
```

```
census_data = census_data.loc[census_data['CHARACTERISTIC_ID'].isin(var_names.keys())]
```

```
# Add in the variable names
```

```
census_data['varname'] = census_data.apply(lambda row: var_names[row['CHARACTERISTIC_ID']], axis=1)
```

```
# Pivot/transpose
```

```
census_data = census_data.pivot(index='GEO_NAME', columns='varname', values='value')
```

```
# Load the EV data
```

```
ev_data = pd.read_csv("ontario_evs_by_fsa_2023-03-31.csv")
```


```
# Join EV and census data by FSA
census_data = census_data.merge(ev_data, left_on='GEO_NAME', right_on='FSA', how='inner')

# Normalize data for individuals over the population
for var in ppl_vars:
    census_data[var] = census_data[var]/census_data['total_pop']

# Normalize data for homes over dwelling count
for var in dwelling_vars:
    census_data[var] = census_data[var]/census_data['total_dwellings']

# Normalize EV data over population, and calculate per-10,000
census_data['BEV'] = (census_data['BEV']/census_data['total_pop'])*10000
census_data['PHEV'] = (census_data['PHEV']/census_data['total_pop'])*10000
census_data['TotalEV'] = (census_data['TotalEV']/census_data['total_pop'])*10000

# Clean up - reorder the columns
all_vars.insert(0, 'FSA')
all_vars.extend(['BEV','PHEV','TotalEV'])
data = census_data[all_vars]
```



Feature Selection

```
#data = pd.read_csv("newdata.csv")
```

Now that the dataset is assembled, we will examine its structure, and select features from each domain that best correlate with our target variable.


```
df = pd.DataFrame(data)
```


```
df.head()
```

	Unnamed: 0	FSA	total_pop	total_dwellings	pop_density	land_area	num_male	ε
0	0	K0A	111626.0	43737.0	NaN	NaN	0.502258	0.0
1	1	K0B	21020.0	9489.0	NaN	NaN	0.504282	0.0
2	2	K0C	52838.0	22714.0	NaN	NaN	0.498978	0.0
3	3	K0E	39649.0	18426.0	NaN	NaN	0.499887	0.0
4	4	K0G	39862.0	19411.0	NaN	NaN	0.496839	0.0

```
df.corr()
```

<b>income_30k</b>	0.142209	0.018681	0.096864
<b>income_40k</b>	0.239515	-0.068326	0.022472
<b>income_50k</b>	0.207168	-0.133492	-0.052951
<b>income_60k</b>	0.029360	-0.192140	-0.114000
<b>income_70k</b>	-0.099302	-0.157246	-0.091124
<b>income_80k</b>	-0.207523	-0.114700	-0.082131
<b>income_90k</b>	-0.183592	-0.094692	-0.086085
<b>income_100k_up</b>	-0.145841	-0.122640	-0.109712
<b>med_total_household_income</b>	-0.244619	0.062834	-0.132781
<b>med_aftertax_household_income</b>	-0.244828	0.084723	-0.124427
<b>indigenous_ppl</b>	0.336357	-0.183989	-0.108537
<b>nonindigenous_ppl</b>	-0.326467	0.219381	0.124441
<b>home_owner</b>	0.076116	-0.122872	-0.097729
<b>home_renter</b>	0.036450	-0.104940	0.080650
<b>home_gov_or_ind_band</b>	0.182273	-0.068847	-0.035424
<b>condo</b>	-0.074746	0.100393	0.161856
<b>non_condo</b>	-0.047327	-0.015677	-0.172645
<b>avg_rooms_home</b>	-0.078504	-0.027388	-0.181258
<b>home_cost_under_30pct</b>	0.209558	-0.353678	-0.124665
<b>home_cost_over_30pct</b>	-0.084205	0.018516	0.162946
<b>can_citizen_ppl</b>	0.051980	-0.196336	-0.179458
<b>non_citizen_ppl</b>	-0.063621	0.241492	0.201821
<b>non_immigrant_ppl</b>	0.157086	-0.300796	-0.186675
<b>immigrant_ppl</b>	-0.175806	0.324770	0.184186
<b>non_perm_res_ppl</b>	0.003372	0.121215	0.162892
<b>no_move_last_yr</b>	0.010188	0.129105	-0.014712
<b>moved_last_yr</b>	-0.033402	-0.089550	0.043140
<b>no_move_5yrs</b>	0.041239	0.105166	-0.003990
<b>moved_last_5yrs</b>	-0.055528	-0.088912	0.024525
<b>school_no_hs</b>	0.254441	0.034664	0.038218
<b>school_hs</b>	0.150044	-0.025609	-0.014352
<b>school_college</b>	0.162289	-0.112132	-0.098212







```
# To find the columns with the highest correlation to TotalEV, we first need to compute the correlation matrix
correlation_matrix = data.corr()
```

```
# Now, we extract the correlations of all features with respect to 'TotalEV'
total_ev_correlations = correlation_matrix['TotalEV'].sort_values(ascending=False)
```

```
<ipython-input-202-010ad1a82039>:2: FutureWarning:
```



The default value of `numeric_only` in `DataFrame.corr` is deprecated. In a future version, it will default to `False`.

total\_ev\_correlations

home_gov_or_ind_band	-0.100223
total_dwellings	-0.103074
age_55	-0.127937
home_owner	-0.129152
no_move_5yrs	-0.129407
homes_mobile	-0.131052
occup_cat_health	-0.131106
work_lastyear_parttime	-0.131691
num_male	-0.133901
homes_duplex_apartment	-0.134685
commute_same_div	-0.134852
occup_ind_agriculture_forestry	-0.147151
can_citizen_ppl	-0.154739
occup_ind_mine_og	-0.161794
homes_lowrise_apartment	-0.161830
age_60	-0.164746
occup_ind_construction	-0.166833
occup_ind_transp_warehs	-0.167163
occup_ind_manuf	-0.168114
no_move_last_yr	-0.170206
homes_other_stationary	-0.170483
commute_time_15	-0.171964
non_condo	-0.173932
edu_field_agriculture_res_env	-0.187902
non_immigrant_ppl	-0.187936
occup_cat_natural_resources_agriculture	-0.193435
occup_ind_other	-0.214784
commute_same_subdiv	-0.241743
income_50k	-0.244129
indigenous_ppl	-0.254281
occup_cat_manuf_util	-0.261751
home_cost_under_30pct	-0.261952
edu_field_health	-0.264745
commute_transp_cardriver	-0.280744
occup_ind_health_socasst	-0.292744
income_10k	-0.296948
work_loc_notfixed	-0.302315
occup_cat_sales_serv	-0.311535
income_20k	-0.313368
occup_ind_accom_food_svc	-0.321089
commute_start_5am	-0.325404
occup_ind_retail_trd	-0.326600
commute_start_7am	-0.331429
commute_time_under15	-0.334368
occup_cat_trades_transp	-0.348162
commute_transp_carpass	-0.377344
income_30k	-0.384509
commute_start_6am	-0.386892
income_40k	-0.399393
edu_field_pers_protect_transp	-0.407856
school_no_hs	-0.431955
school_college	-0.436066
work_loc_workplace	-0.441424
commute_start_noon	-0.464923
school_hs	-0.504329
pop_density	NaN
land_area	NaN

Name: TotalEV, dtype: float64

selected\_features = ["FSA", "med\_fulltime\_income", "income\_100k\_up", "med\_total\_household\_income", "income\_40k",

```
# To find the columns with the highest correlation to TotalEV, we first need to compute the correlation matrix
correlation_matrix = data.corr()

# Now, we extract the correlations of all features with respect to 'TotalEV'
total_ev_correlations = correlation_matrix['TotalEV'].sort_values(ascending=False)

# Next, we'll group the features by their category (e.g., age, commute time, etc.) and select the top
# from each category based on their absolute correlation values with TotalEV.

# Extracting the prefix of each feature to categorize them
feature_categories = total_ev_correlations.index.str.extract(r'([a-zA-Z_]+)').drop_duplicates().dropna()

# Function to get top N features from each category based on absolute correlation with TotalEV
def get_top_n_features_from_each_category(n=3):
    top_features_per_category = {}

    for category in feature_categories[0]:
        # Filtering the columns by category and then taking the top N
        category_features = total_ev_correlations.filter(regex=f'^{category}', axis=0)
        top_features = category_features.abs().sort_values(ascending=False).head(n).index.tolist()
        top_features_per_category[category] = top_features

    return top_features_per_category

top_features_per_category = get_top_n_features_from_each_category()

top_features_per_category
```



```

'occup_ind_health_socasst': ['occup_ind_health_socasst'],
'work_loc_notfixed': ['work_loc_notfixed'],
'occup_cat_sales_serv': ['occup_cat_sales_serv'],
'occup_ind_accom_food_svc': ['occup_ind_accom_food_svc'],
'occup_ind_retail_trd': ['occup_ind_retail_trd'],
'commute_time_under': ['commute_time_under15'],
'occup_cat_trades_transp': ['occup_cat_trades_transp'],
'commute_transp_carpass': ['commute_transp_carpass'],
'edu_field_pers_protect_transp': ['edu_field_pers_protect_transp'],
'school_no_hs': ['school_no_hs'],
'school_college': ['school_college'],
'work_loc_workplace': ['work_loc_workplace'],
'commute_start_noon': ['commute_start_noon'],
'school_hs': ['school_hs'],
'pop_density': ['pop_density'],
'land_area': ['land_area']]

```

```
df_new = df[selected_features]
```

```

# To find the columns with the highest correlation to TotalEV, we first need to compute the correlation matrix
correlation_matrix = df_new.corr()

```

```

# Now, we extract the correlations of all features with respect to 'TotalEV'
total_ev_correlations = correlation_matrix['TotalEV'].sort_values(ascending=False)

```

<ipython-input-207-f27c1cd76d68>:2: FutureWarning:

The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False.

```
total_ev_correlations
```

TotalEV	1.000000
med_fulltime_income	0.671854
income_100k_up	0.462874
occup_cat_snr_mgmt	0.438262
occup_cat_busfin	0.402353
work_loc_home	0.400810
med_total_household_income	0.392518
occup_ind_realestate	0.374130
school_uni_degree	0.373412
edu_field_science	0.352018
edu_field_bus_admin	0.351657
worktype_selfemp	0.342629
condo	0.214618
married_ppl	0.174152
commute_start_9am	0.144460
work_loc_foreign	0.127200
non_citizen_ppl	0.097063
can_citizen_ppl	-0.154739
indigenous_ppl	-0.254281
edu_field_health	-0.264745
work_loc_notfixed	-0.302315
commute_transp_carpass	-0.377344
commute_start_6am	-0.386892
income_40k	-0.399393
edu_field_pers_protect_transp	-0.407856
work_loc_workplace	-0.441424
commute_start_noon	-0.464923
school_hs	-0.504329

Name: TotalEV, dtype: float64

```
df_new.to_csv("data.csv")
```

✓ Data Ingestion Pipeline

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
import plotly.express as px
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.preprocessing import LabelEncoder
```

```
data = pd.read_csv("data.csv")
```

```
pd.set_option('display.max_columns', None)
pd.set_option("display.max_rows", None)
df = pd.DataFrame(data)
df.head(15)
```

	Unnamed: 0	FSA	med_fulltime_income	income_100k_up	med_total_household_incom
0	0	K0A	74500.0	0.127076	115000.
1	1	K0B	56800.0	0.063511	79000.
2	2	K0C	58800.0	0.066429	84000.
3	3	K0E	58000.0	0.063558	83000.
4	4	K0G	64500.0	0.094451	94000.
5	5	K0H	64000.0	0.082900	92000.
6	6	K0J	60000.0	0.071091	76000.
7	7	K0K	58800.0	0.069055	83000.
8	8	K0L	59200.0	0.069489	79000.
9	9	K0M	58800.0	0.074422	80000.
10	10	K1A	NaN	0.075758	73500.
11	11	K1B	65500.0	0.084135	96000.
12	12	K1C	78500.0	0.134417	117000.
13	13	K1E	71000.0	0.104770	110000.
14	14	K1G	68000.0	0.090665	83000.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 518 entries, 0 to 517
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype

```

```
-----
0  Unnamed: 0                518 non-null    int64
1  FSA                      518 non-null    object
2  med_fulltime_income      514 non-null    float64
3  income_100k_up           515 non-null    float64
4  med_total_household_income 515 non-null    float64
5  income_40k               515 non-null    float64
6  occup_cat_snr_mgmt       515 non-null    float64
7  occup_cat_busfin         515 non-null    float64
8  occup_ind_realestate     515 non-null    float64
9  work_loc_home            515 non-null    float64
10 work_loc_workplace       515 non-null    float64
11 work_loc_notfixed        515 non-null    float64
12 school_uni_degree        515 non-null    float64
13 school_hs                515 non-null    float64
14 commute_start_noon       515 non-null    float64
15 edu_field_science       515 non-null    float64
16 edu_field_bus_admin      515 non-null    float64
17 edu_field_health         515 non-null    float64
18 edu_field_pers_protect_transp 515 non-null    float64
19 commute_transp_carpass   515 non-null    float64
20 commute_start_6am        515 non-null    float64
21 commute_start_9am        515 non-null    float64
22 condo                    515 non-null    float64
23 worktype_selfemp         515 non-null    float64
24 indigenous_ppl           515 non-null    float64
25 married_ppl              515 non-null    float64
26 work_loc_foreign         515 non-null    float64
27 can_citizen_ppl          515 non-null    float64
28 non_citizen_ppl          515 non-null    float64
29 TotalEV                  518 non-null    float64
dtypes: float64(28), int64(1), object(1)
memory usage: 121.5+ KB
```

```
df = df.convert_dtypes()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 518 entries, 0 to 517
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Unnamed: 0                            518 non-null    Int64
1   FSA                                  518 non-null    string
2   med_fulltime_income                  514 non-null    Int64
3   income_100k_up                       515 non-null    Float64
4   med_total_household_income           515 non-null    Int64
5   income_40k                           515 non-null    Float64
6   occup_cat_snr_mgmt                   515 non-null    Float64
7   occup_cat_busfin                     515 non-null    Float64
8   occup_ind_realestate                 515 non-null    Float64
9   work_loc_home                        515 non-null    Float64
10  work_loc_workplace                   515 non-null    Float64
11  work_loc_notfixed                    515 non-null    Float64
12  school_uni_degree                    515 non-null    Float64
13  school_hs                            515 non-null    Float64
14  commute_start_noon                   515 non-null    Float64
15  edu_field_science                   515 non-null    Float64
16  edu_field_bus_admin                  515 non-null    Float64
17  edu_field_health                     515 non-null    Float64
18  edu_field_pers_protect_transp         515 non-null    Float64
19  commute_transp_carpass               515 non-null    Float64
20  commute_start_6am                    515 non-null    Float64
21  commute_start_9am                    515 non-null    Float64
22  condo                                515 non-null    Float64
23  worktype_selfemp                     515 non-null    Float64
24  indigenous_ppl                       515 non-null    Float64
25  married_ppl                          515 non-null    Float64
26  work_loc_foreign                     515 non-null    Float64
27  can_citizen_ppl                      515 non-null    Float64
```

```

28 non_citizen_ppl          515 non-null    Float64
29 TotalEV                  518 non-null    Float64
dtypes: Float64(26), Int64(3), string(1)
memory usage: 136.2 KB

```

Each forward sortation area has total number of people residing at that region. Now we have features that describe the total number of people who belong to that particular feature.

## ▼ Missing Values

```
df[df.isna().any(axis=1)]
```

	Unnamed: 0	FSA	med_fulltime_income	income_100k_up	med_total_household_income
<b>10</b>	10	K1A	<NA>	0.075758	7350
<b>158</b>	158	L4V	<NA>	<NA>	<NA>
<b>175</b>	175	L5S	<NA>	<NA>	<NA>
<b>176</b>	176	L5T	<NA>	<NA>	<NA>

```
index_label = [10,158,175,176] #We are removing null values here and the TotalEV bias that you can see
```

```
df = df.drop(index_label)
```

```
df = df.drop('Unnamed: 0', axis=1)
```

```
df.isna().sum()
```

```

FSA                                0
med_fulltime_income                0
income_100k_up                     0
med_total_household_income          0
income_40k                         0
occup_cat_snr_mgmt                  0
occup_cat_busfin                    0
occup_ind_realestate                 0
work_loc_home                       0
work_loc_workplace                  0
work_loc_notfixed                   0
school_uni_degree                   0
school_hs                           0
commute_start_noon                  0
edu_field_science                  0
edu_field_bus_admin                 0
edu_field_health                    0
edu_field_pers_protect_transp        0
commute_transp_carpass               0
commute_start_6am                   0
commute_start_9am                   0
condo                               0
worktype_selfemp                    0
indigenous_ppl                      0
married_ppl                         0
work_loc_foreign                    0
can_citizen_ppl                     0
non_citizen_ppl                     0
TotalEV                             0
dtype: int64

```

df.loc[0]

```
FSA      KOA
med_fulltime_income      74500
income_100k_up      0.127076
med_total_household_income      115000
income_40k      0.078879
occup_cat_snr_mgmt      0.009093
occup_cat_busfin      0.104232
occup_ind_realestate      0.007256
work_loc_home      0.170659
work_loc_workplace      0.267635
work_loc_notfixed      0.075341
school_uni_degree      0.14777
school_hs      0.131869
commute_start_noon      0.038477
edu_field_science      0.013841
edu_field_bus_admin      0.06916
edu_field_health      0.047794
edu_field_pers_protect_transp      0.02598
commute_transp_carpass      0.021276
commute_start_6am      0.092004
commute_start_9am      0.026204
condo      0.026865
worktype_selfemp      0.082911
indigenous_ppl      0.040895
married_ppl      0.554754
work_loc_foreign      0.000806
can_citizen_ppl      0.971234
non_citizen_ppl      0.015677
TotalEV      88.420261
Name: 0, dtype: object
```

df.head()

	FSA	med_fulltime_income	income_100k_up	med_total_household_income	income_40k
0	KOA	74500	0.127076	115000	0.078879
1	KOB	56800	0.063511	79000	0.098004
2	KOC	58800	0.066429	84000	0.096141
3	KOE	58000	0.063558	83000	0.102561
4	KOG	64500	0.094451	94000	0.092561

Continuous Features report

```
def build_continuous_features_report(data_df):
    stats = {
        "Count": len,
        "Miss %": lambda df: df.isna().sum() / len(df) * 100,
        "Card.": lambda df: df.nunique(),
        "Min": lambda df: df.min(),
        "1st Qrt.": lambda df: df.quantile(0.25),
        "Mean": lambda df: df.mean(),
        "Median": lambda df: df.median(),
        "3rd Qrt.": lambda df: df.quantile(0.75),
        "Max": lambda df: df.max(),
        "Std. Dev.": lambda df: df.std(),
    }

    contin_feat_names = data_df.select_dtypes("number").columns
    continuous_data_df = data_df[contin_feat_names]

    report_df = pd.DataFrame(index=contin_feat_names, columns=stats.keys())

    for stat_name, fn in stats.items():
        # NOTE: ignore warnings for empty features
        with warnings.catch_warnings():
            warnings.simplefilter("ignore", category=RuntimeWarning)
            report_df[stat_name] = fn(continuous_data_df)

    return report_df

build_continuous_features_report(df)
```





	Count	Miss %	Card.	Min	1st Qrt.	Me
med_fulltime_income	514	0.0	97	44400.000000	58800.0	68139.1050
income_100k_up	514	0.0	514	0.012731	0.061412	0.1001
med_total_household_income	514	0.0	121	45200.000000	77000.0	94259.1439
income_40k	514	0.0	514	0.047174	0.0734	0.0843
occup_cat_snr_mgmt	514	0.0	507	0.000000	0.003388	0.0072
occup_cat_busfin	514	0.0	514	0.032976	0.069007	0.0913
occup_ind_realestate	514	0.0	513	0.000000	0.006928	0.0104
work_loc_home	514	0.0	514	0.023318	0.079552	0.1326
work_loc_workplace	514	0.0	513	0.167441	0.240323	0.2681
work_loc_notfixed	514	0.0	514	0.017125	0.043569	0.0528
school_uni_degree	514	0.0	514	0.041984	0.10895	0.1897
school_hs	514	0.0	513	0.015504	0.1033	0.1252
commute_start_noon	514	0.0	514	0.011377	0.039829	0.0512
edu_field_science	514	0.0	514	0.002772	0.009196	0.0150
edu_field_bus_admin	514	0.0	514	0.027260	0.056779	0.0786
edu_field_health	514	0.0	514	0.022753	0.044236	0.0508
edu_field_pers_protect_transp	514	0.0	511	0.000000	0.014907	0.0196
commute_transp_carpass	514	0.0	513	0.005306	0.019519	0.0235
commute_start_6am	514	0.0	514	0.000000	0.04431	0.0567
commute_start_9am	514	0.0	514	0.020504	0.035225	0.0430
condo	514	0.0	498	0.000000	0.032001	0.1240
worktype_selfemp	514	0.0	514	0.024189	0.05555	0.0752
indigenous_ppl	514	0.0	510	0.000000	0.009151	0.0365
married_ppl	514	0.0	514	0.319341	0.446095	0.4781
work_loc_foreign	514	0.0	440	0.000000	0.00078	0.0021
can_citizen_ppl	514	0.0	514	0.626602	0.854411	0.8981
non_citizen_ppl	514	0.0	513	0.000000	0.025141	0.0863
TotalEV	514	0.0	514	4.320276	40.721371	78.5657

We see the mean of the TotalEV is 78.56 but ther maximum value is 542.63. We might have to explore on this later

Category Categorical Features Report

```
df.describe(exclude=['number'])
```

	FSA
count	514
unique	514
top	K0A
freq	1

This confirms us that there are no duplicates in the forward sortation area

```
def build_categorical_features_report(data_df):

    def _mode(df):
        return df.apply(lambda ft: ft.mode().to_list())

    def _mode_freq(df):
        return df.apply(lambda ft: ft.value_counts()[ft.mode()].sum())

    def _second_mode(df):
        return df.apply(lambda ft: ft[~ft.isin(ft.mode())].mode().to_list())

    def _second_mode_freq(df):
        return df.apply(
            lambda ft: ft[~ft.isin(ft.mode())]
            .value_counts()[ft[~ft.isin(ft.mode())].mode()]
            .sum()
        )

    stats = {
        "Count": len,
        "Miss %": lambda df: df.isna().sum() / len(df) * 100,
        "Card.": lambda df: df.nunique(),
        "Mode": _mode,
        "Mode Freq": _mode_freq,
        "Mode %": lambda df: _mode_freq(df) / len(df) * 100,
        "2nd Mode": _second_mode,
        "2nd Mode Freq": _second_mode_freq,
        "2nd Mode %": lambda df: _second_mode_freq(df) / len(df) * 100,
    }

    cat_feat_names = data_df.select_dtypes(exclude="number").columns
    continuous_data_df = data_df[cat_feat_names]

    report_df = pd.DataFrame(index=cat_feat_names, columns=stats.keys())

    for stat_name, fn in stats.items():
        # NOTE: ignore warnings for empty features
        with warnings.catch_warnings():
            warnings.simplefilter("ignore", category=RuntimeWarning)
            report_df[stat_name] = fn(continuous_data_df)

    return report_df


build_categorical_features_report(df)
```

	Count	Miss %	Card.	Mode	Mode Freq	Mode %	2nd Mode	2nd Mode Freq	2nd Mode %
FSA	514	0.0	514	NaN	514	100.0	NaN	0	0.0


## Visualizing the Features

As mentioned previously in the continuous feature report section, we will look how the TotalEV feature is dispersed

```
df['TotalEV'].hist(bins=10)
plt.xlabel('TotalEV')
plt.ylabel('Frequency')
plt.title('Histogram of TotalEV')
plt.show()
```



```
plt.bar(df["FSA"], df["TotalEV"])
plt.xlabel('FSA')
plt.ylabel('TotalEV')
plt.title('Barplot of TotalEV')
plt.show()
```




The way the graph has displayed the distribution of values is not bad. The totalEV frequency is decreasing as the number of TotalEV is increasing

```
df.head()
```

	FSA	med_fulltime_income	income_100k_up	med_total_household_income	income_40
0	K0A	74500	0.127076	115000	0.07887
1	K0B	56800	0.063511	79000	0.09800
2	K0C	58800	0.066429	84000	0.09614
3	K0E	58000	0.063558	83000	0.10252
4	K0G	64500	0.094451	94000	0.09256

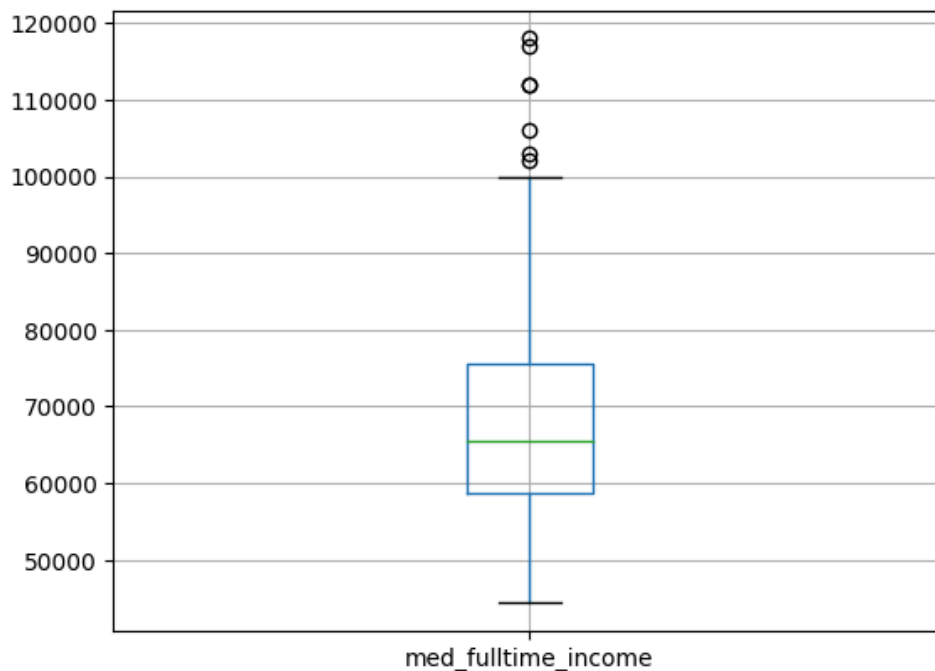
```
df['med_fulltime_income'].hist(bins=20)
plt.xlabel('Median-FT-Income')
plt.ylabel('Frequency')
plt.title('Histogram of Median Total Income')
plt.show()
```



- We see a graphical representation being a little right skewed but not much. The graph is bi-modal, meaning there are two distinct peaks.
- The bars in the right most end have some exceptionally high values. We must investigate if they are outliers or a data entry error

```
df.boxplot(column=['med_fulltime_income'])
```


<Axes: >




Let's try to find out where the median total income is this high. If it's a city or a


```
df.sort_values(by='med_fulltime_income', ascending=False)
```


	FSA	med_fulltime_income	income_100k_up	med_total_household_income	in
307	M5M	118000	0.265738	164000	
41	K2R	117000	0.234438	164000	
330	M8X	112000	0.28238	139000	
291	M4N	112000	0.272138	146000	
285	M4G	106000	0.257424	130000	
185	L6J	103000	0.225635	158000	
293	M4R	102000	0.240574	107000	
19	K1M	100000	0.278681	128000	
295	M4T	100000	0.281165	114000	
284	M4E	98000	0.217289	112000	
45	K2W	98000	0.212206	159000	
297	M4W	97000	0.279158	104000	
43	K2T	97000	0.200489	141000	
327	M6S	97000	0.208205	108000	
30	K2A	97000	0.213919	105000	
273	M3B	96000	0.21118	121000	
303	M5E	96000	0.26785	98000	
106	L1M	95000	0.183657	154000	
188	L6M	95000	0.180435	143000	
168	L5H	94000	0.213683	139000	
23	K1S	93000	0.229069	106000	
289	M4L	93000	0.186852	101000	
187	L6L	93000	0.183427	122000	
50	K4M	92000	0.208623	146000	
51	K4P	92000	0.198124	155000	
439	N7X	92000	0.230292	142000	
305	M5H	92000	0.226293	90000	
42	K2S	91000	0.181606	137000	
129	L2W	91000	0.137825	110000	
481	P1C	91000	0.201721	153000	



































We see that M5M has the highest number of income. It is evident because, M5M represents Toronto, a pretty populous and developed city in terms of number of civilians. Hence we shall not consider this as an outlier

```
#Selecting numeric columns to visualize outliers
numeric_columns = df.select_dtypes(include=["int", "float"]).columns
numeric_columns

Index(['med_fulltime_income', 'income_100k_up', 'med_total_household_income',
      'income_40k', 'occup_cat_snr_mgmt', 'occup_cat_busfin',
      'occup_ind_realestate', 'work_loc_home', 'work_loc_workplace',
      'work_loc_notfixed', 'school_uni_degree', 'school_hs',
      'commute_start_noon', 'edu_field_science', 'edu_field_bus_admin',
      'edu_field_health', 'edu_field_pers_protect_transp',
      'commute_transp_carpass', 'commute_start_6am', 'commute_start_9am',
      'condo', 'worktype_selfemp', 'indigenous_ppl', 'married_ppl',
      'work_loc_foreign', 'can_citizen_ppl', 'non_citizen_ppl', 'TotalEV'],
      dtype='object')


len(numeric_columns)

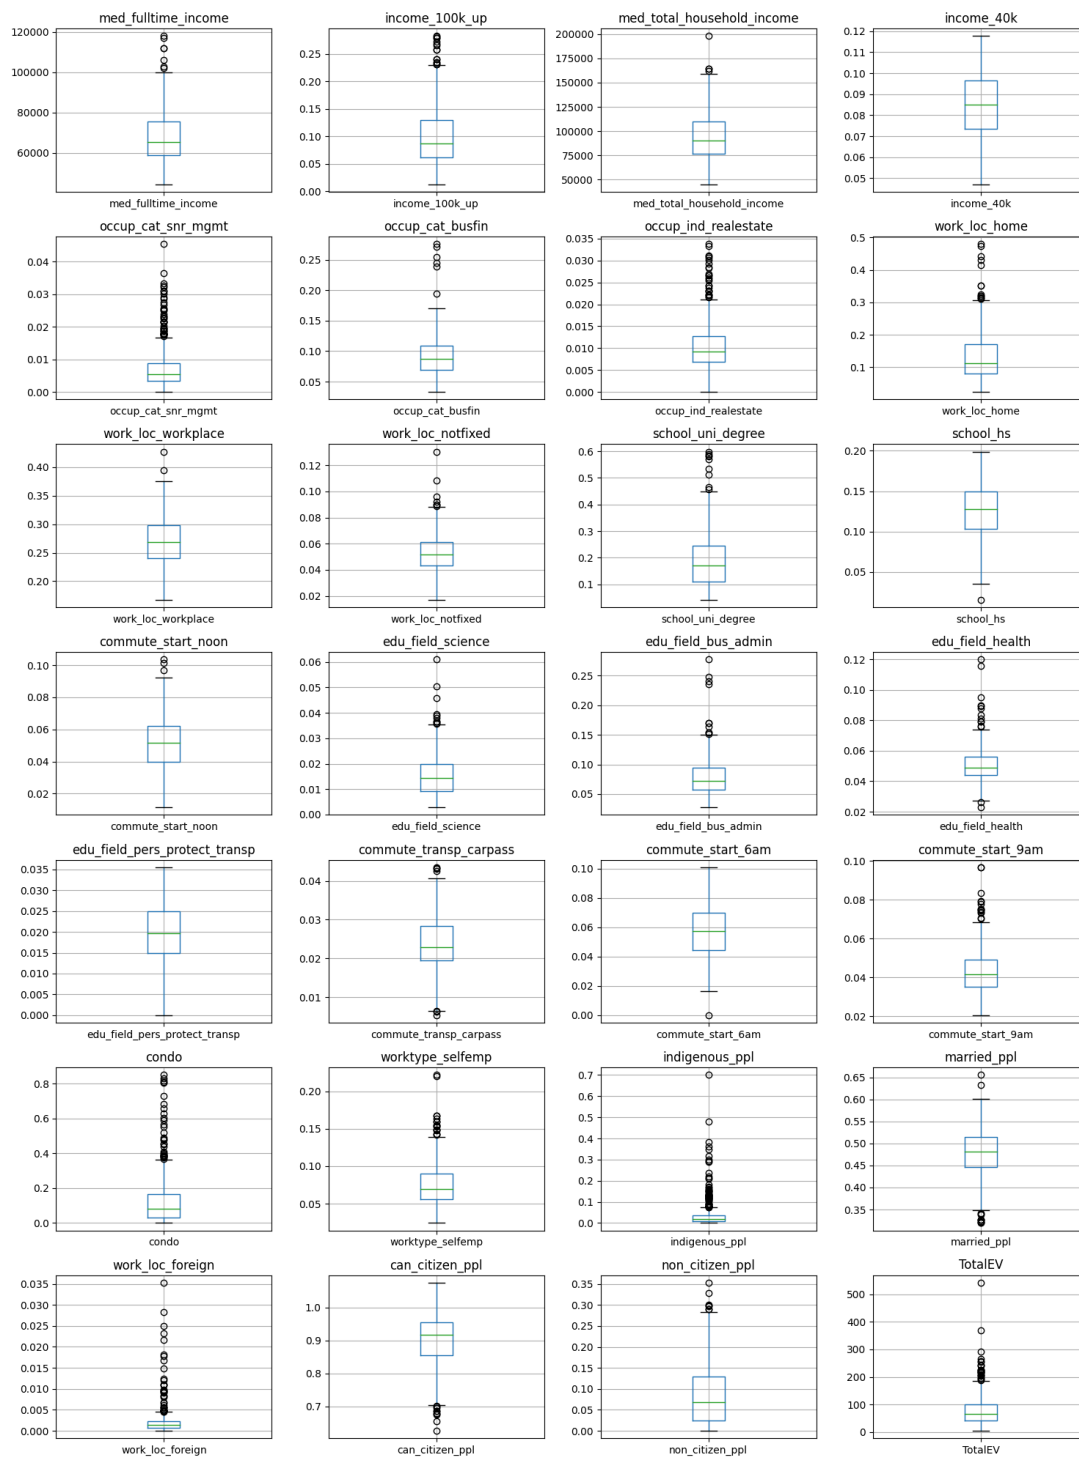
28

fig, axes = plt.subplots(nrows=7, ncols=4, figsize=(15, 20))
axes = axes.flatten()


for i, col in enumerate(numeric_columns):
    ax = axes[i]
    df.boxplot(column=col, ax=ax)
    ax.set_title(col)

plt.tight_layout()
plt.show()
```





```
df['occup_cat_snr_mgmt'].hist(bins=20)
plt.xlabel('Occupation Senior Management')
plt.ylabel('Frequency')
plt.title('Histogram of Senior Management occupation')
plt.show()
```



```
# Step 1: Determine the Bin Ranges
last_three_values = sorted(df['occup_cat_snr_mgmt'].unique())[-3:]
print(last_three_values)

# Step 2: Create a Histogram and store the bin information
n, bins, patches = plt.hist(df['occup_cat_snr_mgmt'], bins='auto') # 'auto' can be replaced with specif:

# Step 3: Annotate the Histogram
# Find the center of each of the last five bins to place the text
for value in last_three_values:
    # Assuming that the value aligns exactly with a bin edge; otherwise, find the nearest bin
    bin_index = (np.abs(bins - value)).argmin() - 1 # Get the index of the bin to the left of
    bin_center = (bins[bin_index] + bins[bin_index + 1]) / 2
    bin_height = n[bin_index]


    # Get the corresponding FSA names for this bin
    fsa_names = df[df['occup_cat_snr_mgmt'] == value]['FSA'].unique()
    fsa_names_str = ", ".join(fsa_names) # Concatenate FSA names into a single string

    # Annotate the histogram with FSA names
    plt.text(bin_center, bin_height, fsa_names_str, ha='center', va='bottom', fontsize=8, rotation=45)

plt.title('Histogram of Senior Management Occupation with FSA Annotations')
plt.xlabel('Occupation Senior Management')
plt.ylabel('Frequency')
plt.show()
```


[0.033342050209205, 0.0364004044489383, 0.0455062571103526]

Histogram of Senior Management Occupation with FSA Annotations



- Distribution Shape:** The distribution of the 'Occupation Senior Management' variable appears to be right-skewed, with a majority of the data concentrated on the left side of the histogram.
- Common Values:** The highest frequency of Senior Management Occupation values is in the lower range (close to 0), indicating that higher percentages in this category are less common.
- Tail of the Distribution:** There is a long tail extending towards the right, which indicates that there are relatively few FSAs with a higher proportion of individuals in senior management roles.
- FSA Annotations:** Annotations on the histogram indicate specific FSAs associated with the higher end of the distribution. These FSAs are likely those with the highest proportion of individuals in senior management occupations, and they stand out from the general trend.
- Outliers:** The presence of FSAs with a higher proportion at the tail end could also indicate potential outliers, or it could simply reflect areas with an unusually high concentration of senior management occupations.
- Analysis Focus:** The annotations indicate specific areas of interest that may warrant further investigation. For example, why do these FSAs have a higher proportion of senior management occupations? This could be related to various factors like the presence of corporate headquarters, higher economic activity, or a different industrial mix compared to other FSAs.

```
df['can_citizen_ppl'].hist(bins=20)
plt.xlabel('Occupation Senior Management')
plt.ylabel('Frequency')
plt.title('Histogram of Senior Management occupation')
plt.show()
```



## ✓ Categorizing the TotalEV using Quantiles and Clustering Similar Features with TotalEV

```
# Calculate the percentiles
#low_threshold = df['TotalEV'].quantile(0.33)
#high_threshold = df['TotalEV'].quantile(0.66)
low_threshold = 50
high_threshold = 100

# Create the categories
df['TotalEV_Category'] = pd.cut(df['TotalEV'], bins=[-float('inf'), low_threshold, high_threshold, float('inf')])


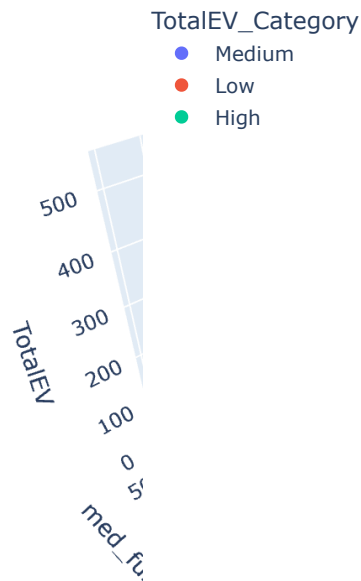
df['TotalEV_Category'].unique()

['Medium', 'Low', 'High']
Categories (3, object): ['Low' < 'Medium' < 'High']

df.head()
```

	FSA	med_fulltime_income	income_100k_up	med_total_household_income	income_40k
0	K0A	74500	0.127076	115000	0.07881
1	K0B	56800	0.063511	79000	0.09800
2	K0C	58800	0.066429	84000	0.09614
3	K0E	58000	0.063558	83000	0.10251
4	K0G	64500	0.094451	94000	0.09251

```
fig = px.scatter_3d(df, z='TotalEV', y='med_fulltime_income', x='med_total_household_income', color='TotalEV_Category')
fig.show()
```




- **Correlation:** There seems to be a positive correlation between `med_total_household_income` and `TotalEV`, as well as between `med_fulltime_income` and `TotalEV`. This suggests that as median incomes increase, the `TotalEV` also tends to increase.
- **Income Disparity:** The spread along the `med_fulltime_income` axis appears to be less pronounced than along the `med_total_household_income`, which might indicate less variability in full-time income as compared to household income.
- **Data Density:** The data points are more concentrated in the lower to middle income brackets for both `med_total_household_income` and `med_fulltime_income`, which suggests that there are fewer data points (FSAs or regions) with very high incomes.
- **Outliers:** There are some outliers, particularly visible in the `TotalEV` dimension (points that are much higher than the main cluster of data), indicating some FSAs or regions with significantly higher values of `TotalEV`.

```
sns.scatterplot(x=df["condo"], y=df["TotalEV"], hue=df["TotalEV_Category"])
```



```
<Axes: xlabel='condo', ylabel='TotalEV'>
```



From this graph, several inferences can be made:

- Positive Correlation:** There appears to be a positive correlation between the 'condo' metric and 'TotalEV'. As the 'condo' metric increases, 'TotalEV' also tends to increase.
- Category Distribution:** The 'TotalEV\_Category' shows that most of the low and medium categories are clustered at the lower end of the 'condo' metric. In contrast, the high category is more spread out and appears more frequently at higher 'condo' values.
- Concentration of Categories:** There's a higher concentration of low and medium 'TotalEV' categories in the lower range of the 'condo' metric. The high 'TotalEV' category seems to be more scattered and less dense, which might suggest that high 'TotalEV' values occur less frequently or are less tightly correlated with the 'condo' metric.
- Data Density:** The majority of the data points are clustered in the lower left corner of the plot, indicating a high density of lower 'TotalEV' values associated with lower 'condo' metric values.
- Outliers:** A few data points in the high 'TotalEV' category stand out significantly above the general trend. These could be outliers or indicate specific instances where the 'TotalEV' is exceptionally high relative to the 'condo' metric.
- Potential Ceiling Effect:** There might be a ceiling effect visible for the 'condo' metric, as there are a number of high 'TotalEV' points that line up along specific 'condo' values.

```
sns.scatterplot(x=df["work_loc_home"], y=df["TotalEV"], hue=df['TotalEV_Category'])
```